

Rater Training Rubric

Hope Programme Comment Generation — Workshop

Why this rubric matters

You will rate every candidate reply on five 1–5 Likert items. Two raters who use these items differently will produce ratings that look like noise to the analysis. The rubric below is what every rater agrees to before the session starts; it is what gives Krippendorff's α a chance.

The five Likert items

Empathy (1–5). Does the reply acknowledge what the participant said and validate the feeling behind it? An empathetic reply names what the participant brought to the post (a goal, a frustration, a small win) before responding.

- **1** — ignores or talks past the participant's content.
- **3** — generic positive acknowledgement (“well done!”) without engaging with what was said.
- **5** — specifically names the feeling or content from the post and validates it in a way that reads as understanding.

Personalisation (1–5). Could this reply be sent to any participant on any post, or does it read as a response to *this* participant in *this* moment?

- **1** — a one-size-fits-all template that would suit any post.
- **3** — mentions the activity type but otherwise generic.
- **5** — references something specific the participant wrote (a name, a goal, a barrier, a small detail) in a way that could not apply to a different post.

Programme fidelity (1–5). Does the reply read as a Hope Programme facilitator's voice, following Motivational Interviewing principles? Specifically: does it avoid prescriptive advice, avoid diagnosis, avoid medication suggestions, and keep the participant's autonomy?

- **1** — breaks programme norms (advice-giving, diagnostic language, prescriptive “you should. . .”).
- **3** — does not break norms, but feels generic rather than programme-aligned.
- **5** — recognisably in the programme's voice (reflective, curious, autonomy-preserving).

Conversational fit (1–5). Would this reply feel natural as the next turn in a peer-support conversation, or does it feel out of place?

- **1** — jarring, formal in a stilted way, or otherwise out of register.
- **3** — works as a reply but does not invite further conversation.
- **5** — conversational, warm, and leaves the door open for the participant to respond.

Length and tone (1–5). Is the reply concise without being curt? Warm without being saccharine? Hope facilitator replies are typically 1–3 sentences. Penalise both very long replies and replies that read as cold.

- **1** — much too long, much too short, or wrong register.
- **3** — acceptable length and tone, nothing notable.
- **5** — exactly the right length and warmth for this post.

The would-send categorical

After the five Likert items, you decide whether you would send this reply. Use these definitions consistently:

- **Send as is** — you would put your name on this reply with no edits.
- **Send with minor edits** — one or two small wording changes (a name correction, a softer hedge, a removed emoji) and you would send it.
- **Send with major edits** — the reply has the right idea but needs substantive rewriting before you would send it.
- **Reject** — you would not send this reply at all; the underlying response is wrong, unsafe, or off-key in a way that editing cannot fix.

Two distinctions worth holding onto

Programme fidelity vs. MI policy compliance. We separately compute a rule-based MI compliance score (counts of prescriptive language, open-ended questions, reflections). Your *programme fidelity* rating is not a redundant rating of the same thing — you should rate whether the *voice* reads as the programme's voice, not whether the count of MI moves is right. A reply can be MI-compliant on paper and still feel off-voice.

Personalisation vs. empathy. Empathy is about whether you acknowledge what the participant brought *at all*. Personalisation is about whether the acknowledgement is specific to this participant. A reply can be empathetic and generic ("that sounds difficult — you've got this") or personalised but cold ("you have been working on your walking goal for three weeks now"). Rate them separately.

When in doubt

Use the middle of the scale (3) rather than the extremes. We will look at the spread across raters, so a deliberate "I am not sure" 3 is informative; a forced 5 is noise.

If you are uncertain whether to flag something as a *Reject* versus *Major edits*, ask yourself: *is the underlying intent of this reply something I would want to land on a participant's post?* If yes, major edits. If no, reject.